

Monolingual HPLT Encoders for the Automatic Evaluation of Answers to Open Questions

Roman Tomchik

Faculty of Mathematics and Physics
Charles University

Alina Haitota

Faculty of Mathematics and Physics
Charles University

Abstract

We describe our submission to the Sensemaking at CLEF 2026 shared task, where a system has to score a student answer to an open question, either on a 0–4 scale or against a rubric, in 13 languages. We use monolingual HPLT encoders and reduce every method to reading one label digit at one position of the prompt: a zero-shot masked-LM baseline, a finetuned classification head, and single-token instruction tuning. Only instruction tuning works, and we argue that this is because it is the only method that reuses the pretrained output head instead of learning a new one. We also build a domain-balanced split, because the released one measures domain transfer rather than grading, and we translate the English rubric into the language of the item, which turns out to be the cheapest improvement we found. Our best macro-domain QWK is 0.334 on the rubric track and 0.301 on the simple track.

1 Task description

This report describes the results of “Sensemaking at CLEF 2026: Automatic Evaluation of Answers to Open Questions” shared task.¹ This task focuses on the “Evaluation” step - score the students’ responses given the context of the topic, question and answer of the student. Some of the items are created by humans, some are LLM generated and checked by humans, and some are fully generated by LLM.

LLMs are very fluent nowadays and seem to be able to solve numerous tasks, but still, they are easy to fool and can give unreliable evaluation. This task tries to get the answer to the question: Can LLMs make sense of the text that they are receiving? Can they be used as reliable judges? Are smaller models being as capable for this task as bigger models?

¹All code, configurations and predictions are available at <https://github.com/welldefinedkernel/sensemaking2026>.

There are two tracks for this task: Simple Rating and Rubric-Based Rating. The Simple Rating track format is: given a triple of texts (context, question, answer), score the answer with the appropriate rating label of 0-4. The Rubric-Based Rating track is: given a quadruple of texts (context, question, rating rubric, answer), label the answer with either No Credit (NC, label 0), Partial Credit (PC, label 1), or Full Credit (FC, label 2), according to the rubric. The rating rubric (three plain textual descriptions) will consist of three parts, describing when each of the possible labels should be applied.

Another important property of this task - it is multilingual. There are 13 languages: English, Czech, Portuguese, German, Hungarian, Finnish, Danish, Swedish, Serbian, Greek, Irish, Romanian, Ukrainian. Some texts are original; some are machine translations of an original in another language.

2 Metrics

Main metrics used are Accuracy and Quadratic Weighted Kappa (QWK; [Cohen, 1968](#)). Accuracy is the fraction of correct labels. It treats all errors as equal and is dominated by the majority class. Always predicting the majority class skews the results. To solve this problem, the metric used to compare models will be the QWK. QWK measures the agreement with the annotator corrected for chance, penalising errors by squared distance. It is important that a constant predictor scores exactly 0, no matter how good its accuracy, which means that QWK cannot be tricked by guessing the majority class. This metric rewards the solution where the ordering of labels is right even if the exact digit at a specific datapoint is wrong.

Each document in the dataset has a domain attached to it. The domains are provided in Table 7. Therefore, the official metric for the task is Macro-domain QWK, which is computed by first com-

puting QWK separately for each domain and then averaging over all the domains. All metrics are computed with scikit-learn (Pedregosa et al., 2011).

3 Dataset analysis

To evaluate models we used trainset and devset located in the shared task GitHub repository, since they contained the correct answers. For this report, this dataset is called regular. There are separate files for two tracks. They contain exactly the same rows and the same ids - only the labels are changed. The trainset has 13856 rows grouped into 2308 questions, the devset has 4146 rows and 739 questions.

All 13 languages are present in both files and have approximately equal frequency (about 1020 training rows and 320 development rows per language), while English is larger (1605 and 344).

The answers are short - the mean is 28 words in devset to 30 words in trainset - but the contexts are larger, and they differ a lot between the two files. In the trainset the mean is 182 while in the devset the mean is 654. On the rubric track, the three criteria description add another 188 words on average and up to 485 to every prompt. This is very important for our choice of the models in the future experiments. Models that have only 512-token window see a truncated prompt on most inputs in the devset, and on the rubric track it runs out of tokens before the context starts at all, so they cannot get the correct score for those inputs.

Another important property of the splits is their domain composition. The trainset contains only two domains, ukr-biology (11384 rows, 82%) and world-history (2472 rows, 18%). The devset contains five: demagog (1770), popular_science (1476), popular (325), world-history (312) and OECD_public (263). Only world-history appears in both splits, and it is 7.5% of the devset, so 92.5% of the development rows (3834 of 4146) come from a domain that the model has not seen during training. So, the smaller model that was taught to grade answers using this trainset will inevitably underperform on devset.

As for the label distributions, on the simple track the trainset is 21.9 / 7.9 / 12.8 / 16.9 / 40.4 percent for labels 0 to 4, and the devset is 7.7 / 16.7 / 26.9 / 5.8 / 42.8. Inside the domains it is worse: training world-history is 56.4% label 0 and only 0.7% label 4, while training ukr-biology is 49.1% label 4. The two training domains have different

label distribution, and so the model taught on those statistics again would transfer badly to devset.

Because of this, we built another domain-balanced split, which we call the balanced dataset. It connects both devset and trainset files, caps the largest ukr-biology domain to twice the size of the next largest domain (from 11384 down to 4500 rows), and then splits every domain and language 80/20. The result has 9747 training rows and 2465 development rows, all six domains present on both sides, matching label distribution - train 20.6 / 11.1 / 17.1 / 14.1 / 37.0 and dev 20.1 / 10.5 / 18.1 / 13.8 / 37.5. Also, since some questions have several examples of answers, the decision was made to have no questions shared between the two splits, otherwise the model would be able to see how the answers should look like for the exact question. Instead, the model can benefit from different types of questions and the correlation between wordings and labels.

4 Architectures used

HPLT models are monolingual encoders, trained on the HPLT web corpora for that language (Burchell et al., 2025; Oepen et al., 2025). Two families were used here. HPLT 2.0 BERT follows the LTG-BERT architecture (Samuel et al., 2023) and has a 512-position window and a masked language modelling head. HPLT 3.0 GPT-BERT (Charpentier and Samuel, 2024) is a hybrid trained with both a masked and a causal objective and has a 16384-position window. The same weights can be used as a masked LM, autoregressive LM, or classification head. During one run we inference all 13 models (for 13 languages) - the input is grouped by language and each group is processed with its own model.

All methods in general have the same approach: reading one digit at one position. A prompt is constructed from the fields of the data point and ends with "The chosen label is: " followed by the slot to be filled, and the system has to put a label digit there. The fixed parts of the prompt - the field labels, the task instruction and the worked examples (for each label the example of the answer that would correspond to this score) where they are used - are machine translated into the language of the item. So a Czech model gets a Czech prompt. To put it simply, the field values are used as they are, since they are already in the target language, everything else in the prompt is translated into the

target language. When a prompt does not fit into the context window, we truncate the context, so the instruction, question, answer and the digit slot will always be given to the model.

We tried three ways to get that digit. The first, `mlm`, is used as a baseline with no training at all: the prompt ends in the model’s mask token, the masked-LM head gives a distribution over the whole vocabulary at that position and from this distribution we take only the logits for the tokens that represent the valid scores, followed by a softmax operation to convert them to probability distribution, and finally argmax operation to get the predicted score. This run does not have the training part, it is only evaluated on the devset.

The second one is `cls`: a sequence classification head is put over the [CLS] representation (Devlin et al., 2019) and the whole model is then finetuned on correct labels, with Cross Entropy Loss used as a criterion. This method requires training, because the head is randomly initialised.

The third, `instruction_tune` - GPT-BERT is loaded as a causal LM, the gold digit is appended after the prompt during training and it is the only supervised position, with every other position label being ignored. At inference the model generates a single token which is later decoded into a digit. Like `cls` it is finetuned on labels, but like `mlm` it reuses the pretrained output head and does not introduce new parameters.

Also, we used two prompt variants. Both have the same instruction, a description of what each score means. The prompt with examples prepends one demonstration per label (exemplary context and question, different answer and the label that scores the answer correctly) before the real item. This is a few-shot prompting tactic (Brown et al., 2020). Detailed prompts can be viewed in the Appendix A.

On top of that there are the two data splits (simple, regular) described above, so that the same model, method and prompt is trained and evaluated twice, once on the original split and once on the balanced one to observe the improvement when the training data is more relevant for the evaluation.

This gives a grid of track x method x prompt x split x model family, and 28 configurations were run. HPLT 2.0 is used only on the simple track, because the rubric prompt with its 190 words of rubric text plus the instruction already overruns 512 tokens, so there is no space for context left. For the same reason HPLT 2.0 is used only with

System	Acc.	QWK _{macro}
3.0 <code>instruction_tune</code> zero_shot	0.533	0.301
3.0 <code>instruction_tune</code> examples	0.521	0.291
3.0 <code>mlm</code> examples	0.274	0.023
3.0 <code>mlm</code> zero_shot	0.169	0.017
3.0 <code>cls</code> zero_shot	0.473	0.013
3.0 <code>cls</code> examples	0.464	0.005
2.0 <code>mlm</code> zero_shot	0.132	−0.004
2.0 <code>cls</code> zero_shot	0.454	−0.008
<i>constant (always 4)</i>	<i>0.375</i>	<i>0.000</i>

Table 1: Simple track (labels 0–4), balanced development set. Rows are sorted by macro-domain QWK.

System	Acc.	QWK _{macro}
3.0 <code>instruction_tune</code> zero_shot	0.132	0.210
3.0 <code>instruction_tune</code> examples	0.110	0.153
3.0 <code>mlm</code> examples	0.206	0.014
3.0 <code>cls</code> zero_shot	0.146	0.012
3.0 <code>mlm</code> zero_shot	0.159	−0.003
2.0 <code>mlm</code> zero_shot	0.193	−0.003
2.0 <code>cls</code> zero_shot	0.309	−0.028
3.0 <code>cls</code> examples	0.224	−0.040
<i>constant (always 4)</i>	<i>0.428</i>	<i>0.000</i>

Table 2: Simple track (labels 0–4), regular development set.

the zero-shot prompt - the few-shot examples do not fit either. Concretely the 28 runs are: 2.0 `mlm` and 2.0 `cls` with zero-shot on the simple track, and 3.0 `mlm`, 3.0 `cls` and 3.0 `instruction_tune` with both prompts on both tracks, each of these on both splits.

The finetuned runs all use the same settings: AdamW optimizer (Loshchilov and Hutter, 2019), learning rate 2e-5, 2 epochs, gradient checkpointing, `max_length` 16384 for 3.0 and 512 for 2.0, and a single generated token for `instruction_tune`. All checkpoints are loaded with the Transformers library (Wolf et al., 2020). The experiments ran on UFAL Artificial Intelligence Cluster using SLURM Job Scheduling system.

5 Results

From the results we can see that only `instruction_tune` demonstrates non-trivial QWK. (Tables 1–4). It is the only method that has macro-domain QWK above zero (between 0.21 and 0.33 in all four combinations of track and split), while everything else is between −0.05 and +0.04. We think that the difference between `cls` and `instruction_tune` models is the output layer: `instruction_tune` reuses the pretrained LM head, `cls` starts from a randomly initialised

System	Acc.	QWK _{macro}
3.0 instruction_tune zero_shot	0.596	0.334
3.0 instruction_tune examples	0.588	0.313
3.0 mlm zero_shot	0.321	0.042
3.0 cls zero_shot	0.483	0.012
3.0 mlm examples	0.390	−0.012
3.0 cls examples	0.449	−0.027
<i>constant (always 2)</i>	<i>0.428</i>	<i>0.000</i>

Table 3: Rubric track (labels 0–2), balanced development set.

System	Acc.	QWK _{macro}
3.0 instruction_tune examples	0.459	0.291
3.0 instruction_tune zero_shot	0.461	0.289
3.0 mlm zero_shot	0.329	0.012
3.0 cls zero_shot	0.332	0.009
3.0 mlm examples	0.359	−0.017
3.0 cls examples	0.340	−0.051
<i>constant (always 2)</i>	<i>0.441</i>	<i>0.000</i>

Table 4: Rubric track (labels 0–2), regular development set.

one. We have only 750 to 1100 training examples per language model, and that is not enough to train a new head. So not adding new parameters is good for these models.

We can also see that the accuracy is a less telling metric here, and this is why the official metric is macro-domain QWK. `cls` with the `zero_shot` prompt on simple/balanced reaches accuracy 0.473, which is better than the constant predictor, but its macro-domain QWK is only 0.013 (Table 1), so inside every domain it agrees with the annotator no better than chance. It just learned which domain gets low scores and which gets high ones instead of grading the answers. `instruction_tune` is different: on the simple track its QWK is positive in all six domains, from 0.093 up to 0.653, and on the rubric track in five domains out of six (Table 7). It really understands the difference between answers inside one domain, and this is what the task asks for.

The zero-shot `mlm` baseline does not work at all. Its macro-domain QWK stays between −0.017 and 0.042 in every run (in both families, with both prompts and on both tracks), and accuracy is lower than the constant predictor everywhere (Tables 1–4). The masked-LM head at the digit position most of the time returns the same label, which can be the most frequent digit during the training of the HPLT models. So encoders of this size cannot grade answers without finetuning. `cls` fails in a

Measure	English	Translated
QWK _{macro} , rubric/balanced	0.272	0.334
QWK _{macro} , rubric/regular	0.256	0.289
Accuracy, rubric/regular	0.445	0.461
QWK, Finnish, rubric/balanced	0.174	0.632
QWK, Ukrainian, rubric/balanced	0.314	0.575

Table 5: Effect of translating the rubric, `instruction_tune` with the `zero_shot` prompt. “English” is the earlier run on the released rubric text, “Translated” is the run reported in Tables 3 and 4.

Language	simple/bal.	rubric/bal.
English	0.716	0.599
Portuguese	0.633	0.707
Irish	0.623	0.568
Swedish	0.576	0.318
Hungarian	0.572	0.631
Danish	0.571	0.543
Ukrainian	0.559	0.575
Czech	0.552	0.542
Finnish	0.552	0.632
German	0.463	0.534
Romanian	0.403	0.277
Serbian	0.397	0.235
Greek	0.348	0.460

Table 6: Macro-domain QWK per language, `instruction_tune` with the `zero_shot` prompt, on the balanced split of both tracks.

different way: it gets a good accuracy because it learns the smallest and the largest score, but rarely predicts the middle labels, and it is visible with macro-domain QWK.

Also we noticed that the examples in the prompt mostly do not help. Nine of the twelve pairs are better with the `zero_shot` prompt, the most visible are `cls` on rubric/regular (0.009 against −0.051, Table 4) and `instruction_tune` on simple/regular (0.210 against 0.153, Table 2). This can happen because the topic of the examples (in our case, LLM-generated examples on boiling water) are not particularly helpful in most of domains. There are three exceptions: `mlm` on both simple splits, where the examples stop the model from labeling everything with score 1, and `instruction_tune` on rubric/regular, where 0.291 against 0.289 is really close.

The split has the biggest effect on the results. For `instruction_tune` the macro-domain QWK using the balanced split grows from 0.210 to 0.301 on the simple track and from 0.289 to 0.334 on the rubric track, but accuracy on the simple track grows much more, from 0.132 to 0.533, which is higher than the constant predictor (Tables 1–4).

Domain	simple	rubric
world-history	0.653	0.615
ukr-biology	0.405	0.524
demagog	0.316	0.422
popular_science	0.168	0.301
OECD_public	0.170	-0.074
popular	0.093	0.214

Table 7: QWK per domain, instruction_tune with the zero_shot prompt, on the balanced split of both tracks.

The training data explains it: the original trainset is 82% Ukrainian biology and its world-history part is 56.4% label 0, so the model learned to give low scores, while the devset is 42.8% label 4. This shows that the original split measures domain transfer abilities of the models and not grading skill, which is the reason why we built the balanced dataset split. On the rubric track the growth is smaller, 0.289 to 0.334 and 0.461 to 0.596, because a 3-point scale has less room for error than a 5-point one.

Translating the rubric helped a lot. In the original data the rubric description was written in English for all 13 languages, so a Czech input has a Czech context, question and answer and about 190 words of English in the middle of the prompt. We machine translated the descriptions into the language of the item, and this gave us improvement without changing the model or the training. For instruction_tune with the zero_shot prompt the macro-domain QWK goes from 0.272 to 0.334 on rubric/balanced and from 0.256 to 0.289 on rubric/regular, and accuracy on the original split goes from 0.445 to 0.461, which is bigger than the constant predictor (Table 5). The gain per language is also high: Finnish goes from 0.174 to 0.632 and Ukrainian from 0.314 to 0.575. The simple track has no rubric, so their numbers did not change.

By language, instruction_tune with the zero_shot prompt on simple/balanced gives English 0.716, Portuguese 0.633, Irish 0.623, Swedish 0.576, Hungarian 0.572, Danish 0.571, Ukrainian 0.559, Czech 0.552, Finnish 0.552, German 0.463, Romanian 0.403, Serbian 0.397 and Greek 0.348 (Table 6). English is first, which is expected, because it is the original language of most contexts and the only one where the prompt is not a translation. But surprisingly Irish is third even though the Irish model probably had less data in its own language during training than, say, German. Greek and Serbian models have the worst performance.

For Serbian, the problem most likely stems from the prompt/model mismatch – the model was finetuned on Cyrillic data, while the dataset for this language is in Latin. On rubric/balanced with translated rubric the order changes: Portuguese 0.707, Finnish 0.632, Hungarian 0.631, English 0.599, Ukrainian 0.575, Irish 0.568, Danish 0.543, Czech 0.542, German 0.534, Greek 0.460, Swedish 0.318, Romanian 0.277 and Serbian 0.235. English is no longer first, while Serbian and Romanian stay at the bottom on both tracks.

We can compare the two HPLT families only on the simple track with the zero_shot prompt, and there both are bad as masked LMs and with a classification head (Tables 1 and 2). The 512-token window is the main limitation, because the devset contexts have 654 words on average, so the 2.0 runs often get the truncated text and do not even see the question. And the long context of 3.0 gives us the possibility to do instruction_tune, and instruction_tune is the only method that works.

6 Summary

To sum up, pretrained monolingual encoders can be finetuned to be able to generate labels. The models do not know how to score answers by themselves, which is expected, given that they are originally foundational language models. Our simple method training makes them usable, but only if it can reuse what the model already has. The best macro-domain QWK we got is 0.334 on rubric/balanced and 0.301 on simple/balanced (Tables 3 and 1), which is fair agreement, but is much lower than the best results by other participants. The main lesson is that the easiest improvement was in the input and not in the model: translating the rubric scored better results than any choice of architecture (Table 5).

References

- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901.
- Laurie Burchell, Ona de Gibert, Nikolay Arefyev, Mikko Aulamo, Marta Bañón, Pinzhen Chen, Mariia

- Fedorova, Liane Guillou, Barry Haddow, Jan Hajič, Jindřich Helcl, Erik Henriksson, Mateusz Klimaszewski, Ville Komulainen, Andrey Kutuzov, Joonas Kytöniemi, Veronika Laippala, Petter Mæhlum, Bhavitvya Malik, and 16 others. 2025. [An expanded massive multilingual dataset for high-performance language technologies](#). *Preprint*, arXiv:2503.10267.
- Lucas Georges Gabriel Charpentier and David Samuel. 2024. [GPT or BERT: why not both?](#) In *The 2nd BabyLM Challenge at the 28th Conference on Computational Natural Language Learning*, pages 262–283, Miami, FL, USA. Association for Computational Linguistics.
- Jacob Cohen. 1968. [Weighted kappa: Nominal scale agreement with provision for scaled disagreement or partial credit](#). *Psychological Bulletin*, 70(4):213–220.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Ilya Loshchilov and Frank Hutter. 2019. [Decoupled weight decay regularization](#). In *International Conference on Learning Representations*.
- Stephan Oepen, Nikolay Arefev, Mikko Aulamo, Marta Bañón, Maja Buljan, Laurie Burchell, Lucas Charpentier, Pinzhen Chen, Mariia Fedorova, Ona de Gibert, Barry Haddow, Jan Hajič, Jindřich Helcl, Andrey Kutuzov, Veronika Laippala, Zihao Li, Risto Luukkonen, Bhavitvya Malik, Vladislav Mikhailov, and 13 others. 2025. [HPLT 3.0: Very large-scale multilingual resources for LLM and MT. mono- and bilingual data, multilingual evaluation, and pre-trained models](#). *Preprint*, arXiv:2511.01066.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. 2011. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830.
- David Samuel, Andrey Kutuzov, Lilja Øvrelid, and Erik Velldal. 2023. [Trained on 100 million words and still in shape: BERT meets British National Corpus](#). In *Findings of the Association for Computational Linguistics: EACL 2023*, pages 1954–1974, Dubrovnik, Croatia. Association for Computational Linguistics.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven
- Le Scao, Sylvain Gugger, and 3 others. 2020. [Transformers: State-of-the-art natural language processing](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online. Association for Computational Linguistics.

A Prompt templates

A.1 simple/zero_shot

You grade how well an Answer responds to the Question given the Context.

Choose exactly one label on a 0-4 scale:

0: incorrect or unrelated; contradicts the Context.

1: mostly wrong; only a minor relevant element.

2: partially correct; misses key information or mixes in errors.

3: largely correct; supported by the Context with minor omissions.

4: fully correct and complete; well supported by the Context.

Context: <CONTEXT>

Question: <QUESTION>

Answer: <ANSWER>

The chosen label is: [MASK]

A.2 simple/examples

The same instruction as A.1, followed by five demonstrations that share one short context, and then the real item in the form shown in A.1.

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Answer: Water never boils; it only freezes.

The chosen label is: 0

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Answer: It has something to do with heat.

The chosen label is: 1

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Answer: Around 90 degrees Celsius.

The chosen label is: 2

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Answer: About 100 degrees.

The chosen label is: 3

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Answer: 100 degrees Celsius at sea level.

The chosen label is: 4

A.3 rubric/zero_shot

You grade how well an Answer responds to the Question given the Context, following the Rubric written for this Question.

Choose exactly one label on a 0-2 scale:

0: the Answer meets the No Credit criterion.

1: the Answer meets the Partial Credit criterion.

2: the Answer meets the Full Credit criterion.

Context: <CONTEXT>

Question: <QUESTION>

Rubric:

0 (No Credit): <NC CRITERION>

1 (Partial Credit): <PC CRITERION>

2 (Full Credit): <FC CRITERION>

Answer: <ANSWER>

The chosen label is: [MASK]

A.4 rubric/examples

The same instruction as A.3, followed by three demonstrations that share one context and one rubric, and then the real item in the form shown in A.3.

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Rubric:

0 (No Credit): The answer gives no temperature, or one that contradicts the Context.

1 (Partial Credit): The answer refers to boiling or heat but not to 100 degrees Celsius.

2 (Full Credit): The answer states 100 degrees Celsius.

Answer: Water never boils; it only freezes.

The chosen label is: 0

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Rubric:

0 (No Credit): The answer gives no temperature, or one that contradicts the Context.

1 (Partial Credit): The answer refers to boiling or heat but not to 100 degrees Celsius.

2 (Full Credit): The answer states 100 degrees Celsius.

Answer: It boils once it gets hot enough.

The chosen label is: 1

Context: Water boils at 100 degrees Celsius at sea level.

Question: At what temperature does water boil at sea level?

Rubric:

0 (No Credit): The answer gives no temperature, or one that contradicts the Context.

1 (Partial Credit): The answer refers to boiling or heat but not to 100 degrees Celsius.

2 (Full Credit): The answer states 100 degrees Celsius.

Answer: 100 degrees Celsius at sea level.

The chosen label is: 2

A.5 A translated prompt

The same simple/zero_shot prompt from A.1 as the Czech model actually receives it. Every fixed string has been translated; only the field values are left untouched.

Hodnotíte, jak dobře odpověď odpovídá na otázku v kontextu.

Vyberte přesně jeden štítek na stupnici 0-4:

0: nesprávné nebo nesouvisející; odporuje kontextu.

1: většinou špatně; pouze vedlejší relevantní prvek.

2: částečně správně; postrádá klíčové informace nebo míchá chyby.

3: z velké části správně; podpořeno Kontextem s drobnými opomenutími.

4: zcela správně a úplně; dobře podporované Kontextem.

Kontext: <KONTEXT>

Otázka: <OTÁZKA>

Odpověď: <ODPOVĚĎ>

Zvolený štítek je: [MASK]